

Java企业应用 -性能优化原则, 方法与策略

阿里巴巴 李三红

自我介绍

- 2014年 加入蚂蚁金服，目前在阿里基础架构事业群，基础软件部门：
 - 开发基于OpenJDK阿里定制版本：AJDK
 - 开发性能故障诊断工具：ZProfiler，ZDebugger，PIPA
- 联系方式
 - mail: sanhong.lsh@alibaba-inc.com
 - 微博: [sanhong_li](#)

AJDK



Performance
Analysis

Outline

1. Performance Basics and Methodology
2. Fundamentals of Performance Tuning
 - Profiling Driven Optimization
 - JVM Tuning
 - GC
 - JIT
3. Optimization Strategy for JavaEE
4. Recap

Recall Little's law

$$L = \lambda * W$$

In queueing theory, the long-term **average number of customers** in a stable system, **L**, is equal to the long-term **average arrival rate**, **λ** , multiplied by the **average time a customer spends** in the system, **W**

source: https://en.wikipedia.org/wiki/Little%27s_law

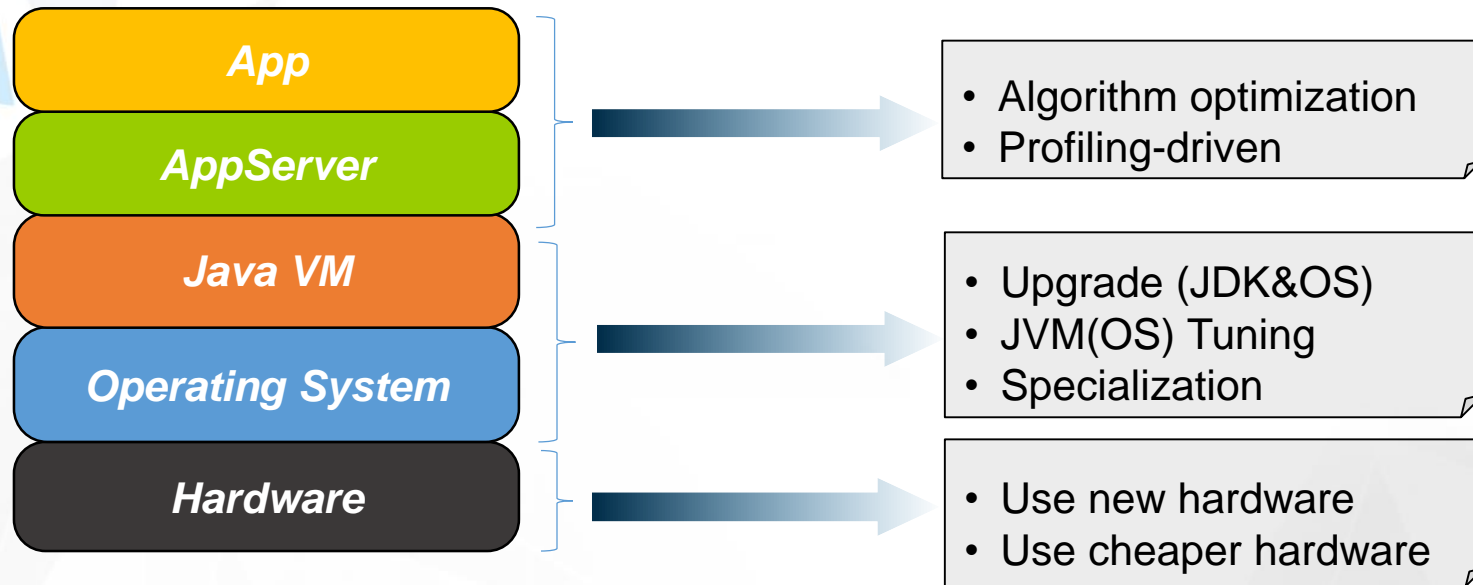
Throughput and RT

$$\text{MeanNumberInSystem} = \text{MeanThroughput} * \text{MeanResponseTime}$$

- Throughput and RT are related
 - ✓ Decreasing RT “almost always” improves Throughput
 - ✓ Throughput improving doesn’t necessarily mean RT decreasing
- Performance tuning and cost saving
 - ✓ More higher throughput/lower RT **but without** adding new hardware

source: https://en.wikipedia.org/wiki/Little%27s_law

Approaches to Performance



Java's view

Approaches:

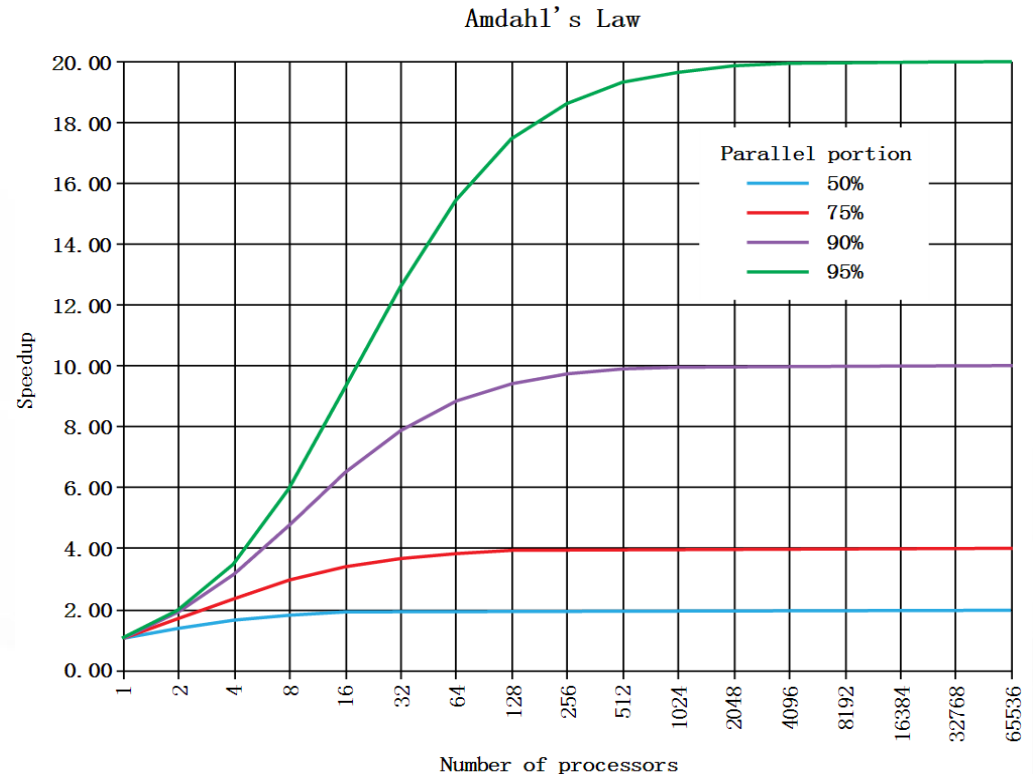
- Outside in approach(performance baseline)
- Layered approach("Bottom up" or "Top down")
- A hybrid of both a), b)

Amdahl's Scaling Law

$$\text{Speedup} = \frac{1}{F + \frac{1-F}{N}}$$

F: is fraction of work that is serial

N: is number of threads



source: https://en.wikipedia.org/wiki/Amdahl%27s_law

- Reduce the amount of serial work performed


Costs Reduce Scaling

1. Potential contributors to **F**:

- Synchronization(**synchronized**&**j.u.c.Lock**)

- data structures need to be thread safe
- communication overhead between threads

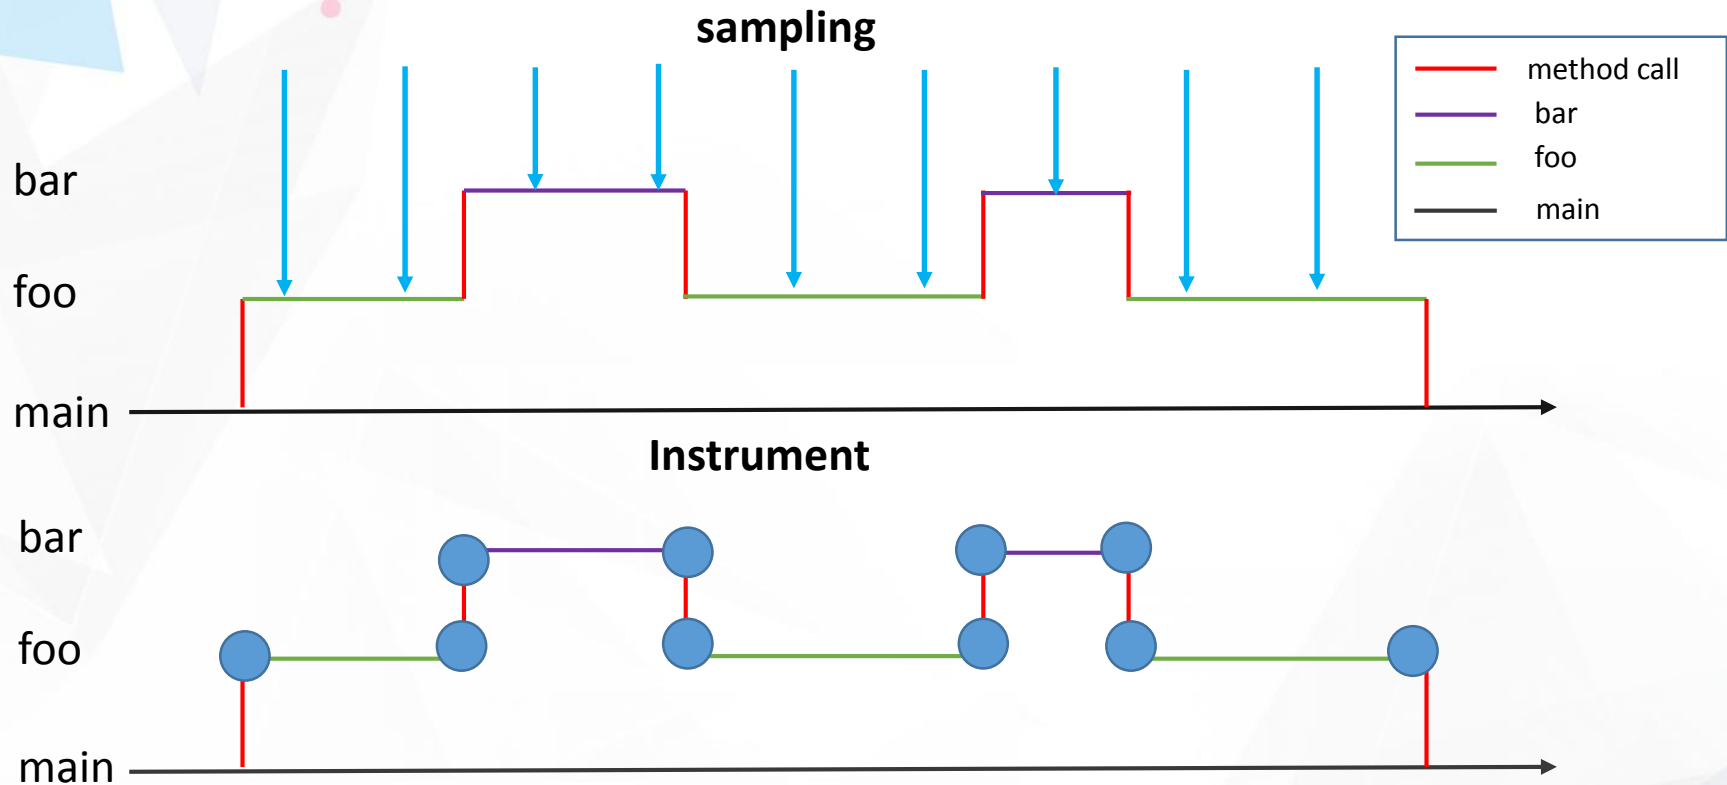
- Infamous “stop the world” (aka STW) in JVM

- 
- JConsole (MXBean)
 - Java Mission Control
 - JProfiler
 - HealthCenter&jucProfiler

2. Cost incurred when the **N** gets increased

- Thread context switch

Profiling: Sampling vs Instrument



Available Technology:

BCI, JVMTi, javax.management, System.currentTimeMillis()

Sampling vs Instrument

■ Sampling

- ✓ Lower overhead (determined by sampling interval)
- ✓ Discover unknown code
- ✓ Non intrusive
- ✓ No execution path
- ✓ Periodicity Bias

■ Instrument

- ✓ Wall time (estimate IO time)
- ✓ Full execution path
- ✓ Configuration on what methods to instrument
- ✓ Generally more data to be collected

Safepoint Bias

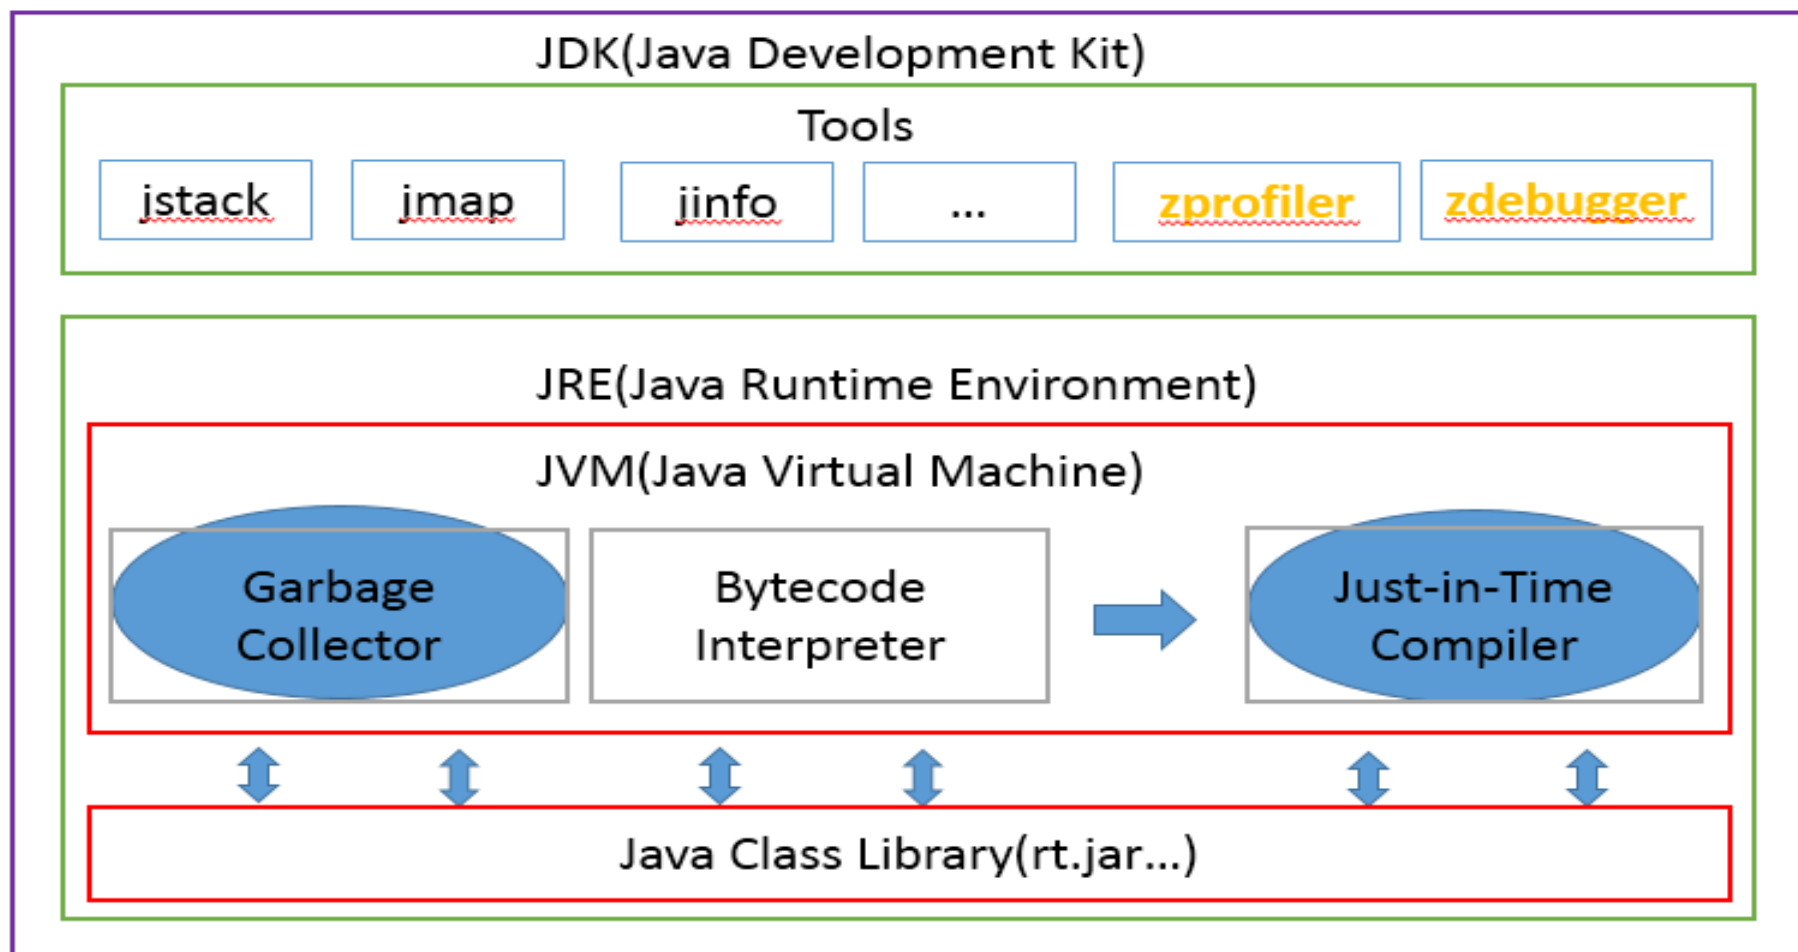
- Stack trace sampling happens only when the given thread at a *safepoint*
 - ✓ The hot loop may not get profiled anymore
- Use following tools instead
 - ✓ Java Mission Control
 - ✓ Honest Profiler(github)
 - ✓ **ZProfiler**(alipay internal profiling tool)

Tools for Diagnostics

JMX **jcmd** attach **jinfo**
jstatd java mission control
java flight recorder **jps** **jstat**
jvisualvm **jstack** perf counters
jmap serviceability agent

- Most of them could be found in JAVA_HOEM/bin
- Good reference: Troubleshooting Guide for JavaSE 6 with HotSpotVM

Basics of JVM Tuning



Guild for GC Tuning

- **Select** the right GC algorithm
 - parallel old ,CMS and G1 collector
 - Rule of thumb: GC overhead is ideally $< 10\%$
- **Choose** the right heap size

Space	Command Line Option	Occupancy Factor
Java heap	-Xms and -Xmx	3x to 4x old generation space occupancy after full garbage collection
Permanent Generation	-XX:PermSize -XX:MaxPermSize	1.2x to 1.5x permanent generation space occupancy after full garbage collection
Young Generation	-Xmn	1x to 1.5x old generation space occupancy after full garbage collection
Old Generation	Implied from overall Java heap size minus the young generation size	2x to 3x old generation space occupancy after full garbage collection

source: Charlie Hunt, Binu John Java™ performance

- **Configure** the appropriate GG parameters

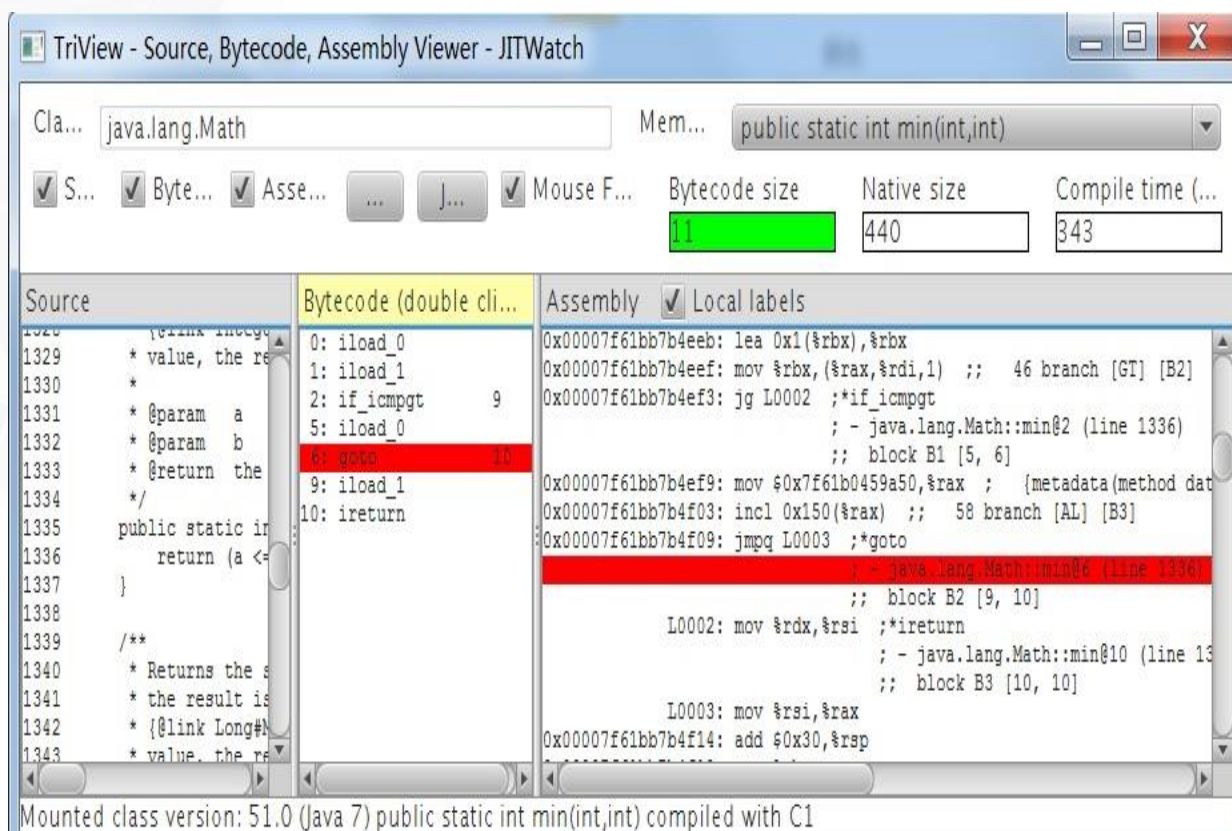
JIT and common optimization

- Important concepts
 - Profiler guided optimization(PGO)
 - Optimization decisions are made dynamically
 - Mix mode execution
- Some common optimization
 - Inlining
 - Intrinsic
 - Monomorphic dispatch
 - ...

☐ **Liskov substitution principle**
Subtypes **MUST** be substitutable
for their base types

JIT Profiling with JITWatch

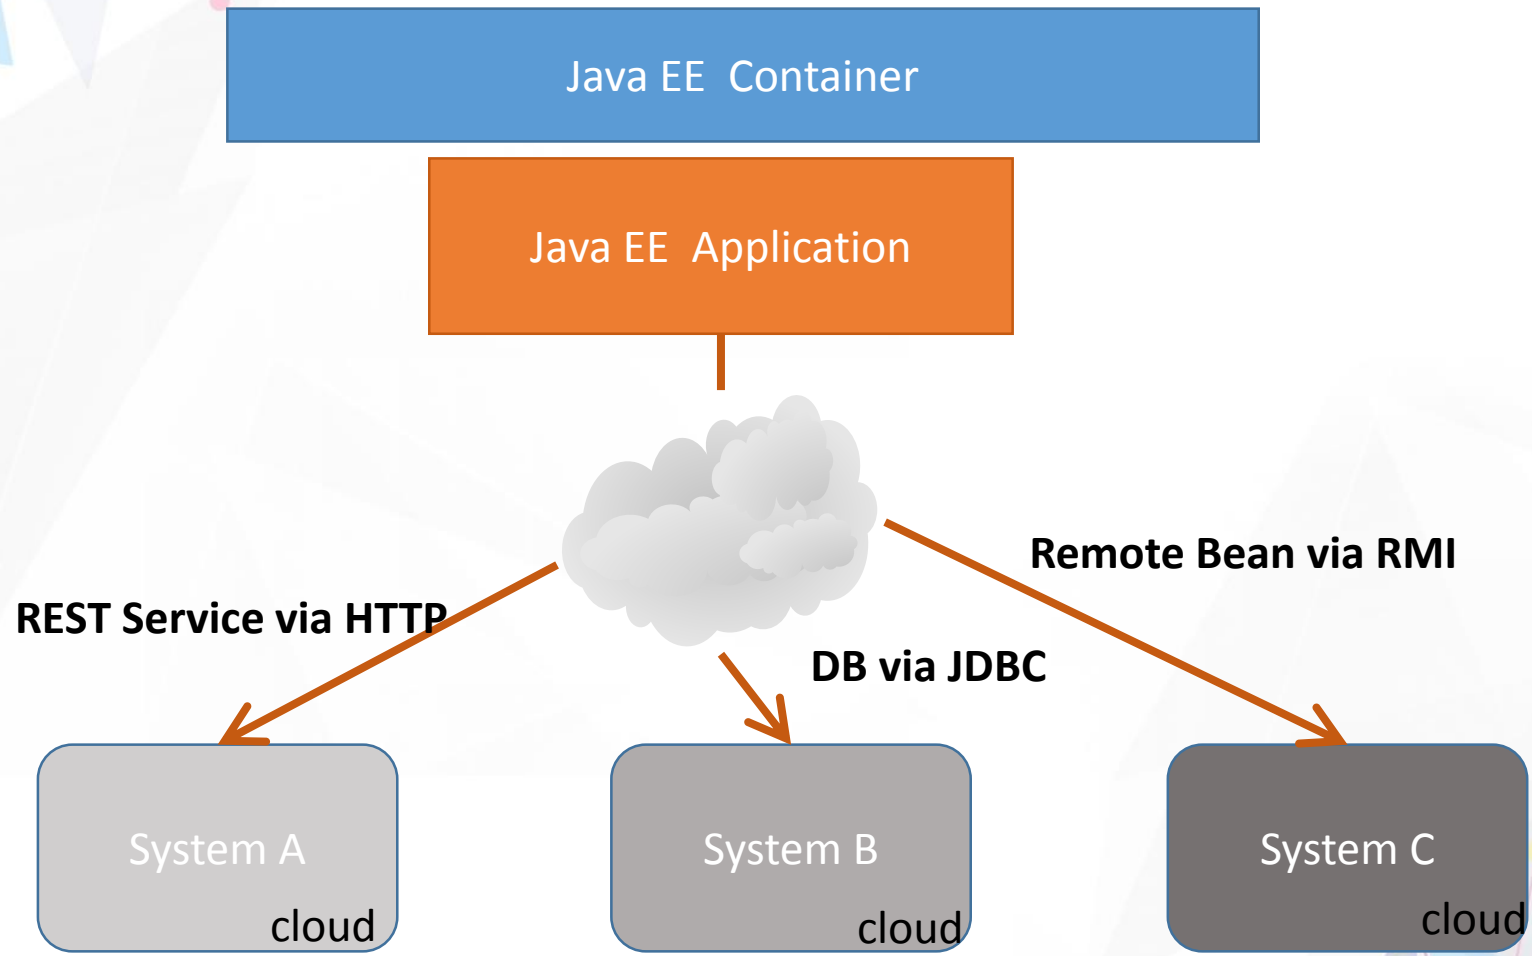
- JITWatch: a **graphical visualization and analysis tool** for understanding the JIT



Enabled by:

- XX:+UnlockDiagnosticVMOptions
- XX:+TraceClassLoading
- XX:+LogCompilation
- XX:+PrintAssembly


Typical distributed JEE architecture



The problem...

- Add communication cost
 - ✓ RPC
 - ✓ serialization/deserialization
- Can **not** shift resources towards demand
- Can **not** share the underlying Java artifacts(such as JIT)

Multitenancy for JavaEE



- ❑ Run multiple Java EE applications (**as tenants**) into same Java EE container

High Density Cloud for JavaEE

The JavaEE applications developed **separately** can be deployed **seamlessly** into the same container.


- devOps
 - ✓ Orchestrate JavaEE application at scale
- Infrastructure
 - ✓ 'Multi-tenant' JavaEE container
 - ✓ 'Virtualized' JVM



source: <https://www.dreamstime.com>

Tomcat/JDK extended for PaaS

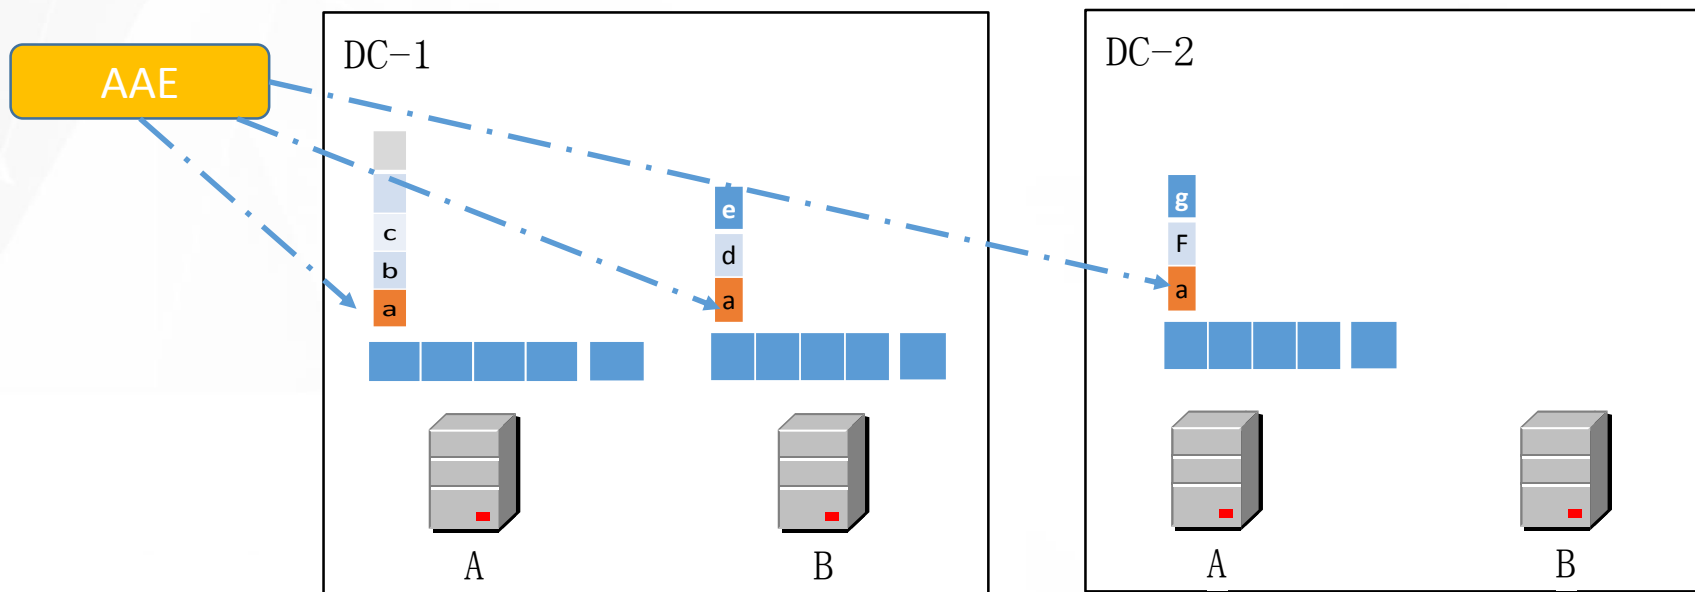
- **AliTomcat**: run multiple apps side-by-side safely
- **AJDK** allows for collocation of multiple JEE apps(as tenant) in a single instance of JVM:
 - **Isolate** application from one another.
 - **Share** metadata aggressively and transparently, such as:
 - ✓ bytecodes of methods
 - ✓ GC
 - ✓ JIT




AJDK : Alibaba/Alipay JDK, based on OpenJDK

AAE: Alibaba Application Engine

- Scaling tenant application with AAE
 - **spread** application evenly across hosts
 - but **pack** applications on the single JVM as much as possible, based on its resource capacity:
 - CPU usage
 - Memory(monitoring GC)




Benefits




Benefits(Cont.)

- **Eliminate** the unnecessary RPC
- **Minimize** the cost caused by object serialization/deserialization
- **Share** underlying Java artifacts as much as possible
 - GC
 - JIT
 - Heap

Compared with Docker



CPU overcommit via docker



'shared memory' via Multitenancy

Summary

- **What we covered:**
 - Performance basics& methodology
 - Performance tuning
 - Profiling
 - Tuning from JVM perspective
 - Multitenancy for JavaEE

THANKS

SequeMedia
盛拓传媒

IT168.com
专注IT 10年

ChinaUnix

ITPUB
www.itpub.net

reach to me:

mail: sanhong.lsh@alibaba-inc.com

weibo: [sanhong_li](#)